

CONTACT Email: nitish@nyu.edu
INFORMATION Website: <https://joshinh.github.io>

EDUCATION **New York University** 2020-Present
Ph.D. Computer Science, Courant Institute
Advisor: Prof. He He
GPA: 4.0/4.0

Indian Institute of Technology (IIT) Bombay 2016-2020
B.Tech. (with Honors) in Computer Science
Advisor: Prof. Preethi Jyothi
GPA: 9.40/10

HONORS AND AWARDS Dean's Dissertation Fellowship, 2024-2025
NSF Graduate Research Fellowship, 2020
Henry M. MacCracken Fellowship, 2020
Institute Academic Award for Academic Excellence at IIT Bombay, 2017
All India Rank 176 in JEE-Advanced for entrance to IITs, 2016
Department Travel Award for ACL 2019 (IIT Bombay)

PUBLICATIONS AND PREPRINTS

- [1] LLMs Are Prone to Fallacies in Causal Inference
Nitish Joshi, Abulhair Saparov, Yixin Wang, He He
Preprint, Under Submission
- [2] Personas as a Way to Model Truthfulness in Language Models
Nitish Joshi*, Javier Rando*, Abulhair Saparov, Najoung Kim, He He
Preprint, Under Submission
- [3] Nuisances Via Negativa: Uncovering Spurious Features via Semantic Corruption
Aahlad Puli, **Nitish Joshi**, Yoav Wald, He He and Rajesh Ranganath
Transactions on Machine Learning Research (TMLR)
- [4] Improving Multi-Hop Reasoning in LLMs by Learning from Rich Human Feedback
Nitish Joshi, Koushik Kalyanaraman, Zhiting Hu, Kumar Chellapilla, He He and Li Erran Li
AAAI Workshop on Neuro-Symbolic Learning and Reasoning in the era of Large Language Models, 2024
- [5] Testing the General Deductive Reasoning Capacity of Large Language Models Using OOD Examples
Abulhair Saparov, Richard Yuanzhe Pang, Vishakh Padmakumar, **Nitish Joshi**, Seyed Mehran Kazemi, Najoung Kim*, He He*
Neural Information Processing Systems (NeurIPS), 2023
- [6] Measuring Inductive Biases of In-Context Learning with Underspecified Demonstrations
Chenglei Si*, Dan Friedman*, **Nitish Joshi**, Shi Feng, Danqi Chen and He He
Association for Computational Linguistics (ACL), 2023
- [7] Are All Spurious Correlations in Natural Language Alike? An Analysis through a Causal Lens
Nitish Joshi*, Xiang Pan* and He He
Empirical Methods in Natural Language Processing (EMNLP), 2022

- [8] An Investigation of the (In)effectiveness of Counterfactually Augmented Data
Nitish Joshi and He He
Association for Computational Linguistics (ACL) 2022 (Part of the work also accepted at ICLR 2021 Robust ML Workshop)
- [9] QuALITY: Question Answering with Long Input Texts, Yes!
 Richard Yuanzhe Pang*, Alicia Parrish*, **Nitish Joshi***, Nikita Nangia, Jason Phang, Angelica Chen, Vishakh Padmakumar, Johnny Ma, Jana Thompson, He He and Sam Bowman
North American Chapter for the Association for Computational Linguistics (NAACL) 2022
- [10] Coupled Training of Sequence-to-Sequence Models for Accented Speech Recognition
 Vinit Unni*, **Nitish Joshi*** and Preethi Jyothi
International Conference on Acoustics, Speech, & Signal Processing (ICASSP), 2020
- [11] Explore, Propose and Assemble: An Interpretable Model for Multi-Hop Reading Comprehension
 Yichen Jiang*, **Nitish Joshi***, Yen-Chun Chen and Mohit Bansal
Association for Computational Linguistics (ACL), 2019
- [12] Cross-Lingual Training for Automatic Question Generation
 Vishwajeet Kumar, **Nitish Joshi**, Arijit Mukherjee, Ganesh Ramakrishnan and Preethi Jyothi
Association for Computational Linguistics (ACL), 2019

RESEARCH AND
INDUSTRY
EXPERIENCE

Google, California

Summer 2024

Research Internship

Advisor: Ankur Taly

Working with the Cloud AI team to improve factuality and context grounding in large language models.

Amazon AI, California

Summer 2022

Applied Scientist Internship

Advisor: Erran Li

Worked with the human-in-the-loop team at Amazon AWS AI, on learning from rich human feedback to improve reasoning in language models.

NYU, Machine Learning for Language Lab

2020-Present

Research Assistant

Advisor: Prof. He He

Broadly working in robustness of Natural Language Understanding models.

IIT Bombay, CSALT Lab

2018-2020

Undergraduate Research Assistant

Advisor: Prof. Preethi Jyothi

Worked on multiple projects including cross-lingual question generation, accented speech recognition, and learning with noisy label (Thesis)

NEC Labs, Princeton

Summer 2019

Research Internship

Advisor: Dr. Christopher Malon

Worked on building explainable models for natural language inference

University of North Carolina, Chapel Hill, NLP Lab

Summer 2018

Research Internship

Advisor: Prof. Mohit Bansal

Worked on building interpretable models for multi-hop reading comprehension (published at ACL 2019)

TEACHING
EXPERIENCE

CS2590: Natural Language Processing

Spring 2023

New York University

Teaching Assistant with Prof. He He. Involved in teaching sections, creating homeworks, advising and grading projects.

NYU AI School

Jan 2021

New York University

Teaching Assistant for a week-long winter school on machine learning for early year undergraduates and underrepresented groups in CS.

CS726: Advanced Machine Learning

Spring 2020

Indian Institute of Technology, Bombay

Teaching Assistant with Prof. Sunita Sarawagi.

SERVICE

Reviewer (conferences): COLM 2024, ACL Rolling Review (2024), EMNLP 2023, ACL 2023, EMNLP 2022

Reviewer (workshops): ICLR 2024 Workshop on Mathematical and Empirical Understanding of Foundation Models, Workshop on Spurious Correlations, Invariance, and Stability, ICML 2022 & ICML 2023, Workshop on Efficient Benchmarking in NLP, ACL 2022.

RELEVANT
COURSES

Advanced Machine Learning, Automatic Speech Recognition, Foundations of Intelligent and Learning Agents, Optimization, Deep Learning.

REFERENCES

Available on Request.